

# REINFORCEMENT LEARNING APPLIED TO MOBILE ROBOT NAVIGATION

---

Navyata Sanghvi

May 14, 2015

**Project Guides:**

Dr. P.V. Manivannan (ME, IIT Madras)

Dr. Balaraman Ravindran (CSE, IIT Madras)

## 1. Preliminaries

- Markov Decision Process (MDP)
- Value Function
- Q-Learning Algorithm

## 2. Robot Navigation Tasks


- Any Goal; No Obstacles
- Fixed Goal; Fixed Obstacle Configuration
- Any Goal; Any Obstacle Configuration

## 3. Conclusion and Future Work

# PRELIMINARIES

---

# MARKOV DECISION PROCESS (MDP)



**Definition 1:** Markov Decision Process (MDP)  $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- $\mathcal{S}$ : state space
- $\mathcal{A}_s$ : action space
- $\mathcal{P}(s'|s, a)$ : probability that an action  $a$  taken in state  $s$  will lead to state  $s'$  at the next time step
- $R_a(s, s')$ : immediate reward received after transition from state  $s$  to  $s'$  upon taking action  $a$
- $\gamma$ : discount factor  $\in (0, 1)$

**Definition 2:** A policy  $\pi$  for an MDP is a mapping  $\pi : \mathcal{S} \mapsto \mathcal{A}$  from states to actions;  $\pi(s)$  denotes the action choice in state  $s$ .

**Definition 3:** The cumulative discounted sum of rewards  $R_\infty$  is as follows:


$$R_\infty \stackrel{\text{def}}{=} \sum_{t=0}^{\infty} \gamma^t R_{a_t}(s_t, s_{t+1}),$$

$$\text{where } a_t = \pi(s_t)$$

**Definition 4:** The optimal policy  $\pi^*$  is as follows:

$$\pi^* \stackrel{\text{def}}{=} \arg \max_{\pi} \mathbb{E}[R_\infty]$$

# MARKOV DECISION PROCESS (MDP) EXAMPLE



1.  $\gamma = 0.1$

- On turning left,  $R_\infty = -20 + 0.1(0) + 0.1^2(100) = -19$ .
- On turning right,  $R_\infty = 0.1^2(20) = 0.2$ .

2.  $\gamma = 0.9$

- On turning left,  $R_\infty = -20 + 0.9(0) + 0.9^2(100) = 61$ .
- On turning right,  $R_\infty = 0.9^2(20) = 16.2$ .

**Definition 5:** A state-action value function is a mapping  $Q : \mathcal{S} \times \mathcal{A} \mapsto \mathbb{R}$ .

$$Q^\pi(s, a) \stackrel{\text{def}}{=} \sum_{s'} P(s'|s, a) \left( R_a(s, s') + \gamma Q^\pi(s', \pi(s')) \right)$$

## Q-Learning

Bellman Operator B

$$BQ(s, a) \stackrel{\text{def}}{=} \sum_{s'} P(s'|s, a) \left( R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$

Under optimal Policy  $\pi^*$

$$Q^{\pi^*}(s', \pi^*(s')) = \max_{a'} Q(s', a')$$

$$\Leftrightarrow BQ^* = Q^*$$

## Updates

$$Q(s, a) \leftarrow \sum_{s'} P(s'|s, a) \left( R_a(s, s') + \gamma \max_{a'} Q(s', a') \right)$$
$$\pi(s) \leftarrow \arg \max_a Q(s, a)$$

---

### Algorithm 1: Q Learning

---


```
1: Input: Bellman operator  $B$  for the MDP, and allowable error  $\epsilon$ 
2: Initialize  $Q(s, a) \leftarrow 0$ 
3: Initialize  $converged = false$ 
4: repeat
5:    $Q_{new} = BQ_{old}$ 
6:   if  $\max(|Q_{new} - Q_{old}|) < \epsilon$  then
7:      $converged = true$ 
8:   end if
9:    $Q_{new} = Q_{old}$ 
10: until  $converged = true$ 
11: return  $Q_{new}$ 
```

---




# ROBOT NAVIGATION TASKS

---

State and Action Spaces

$$a = (m_{\text{change}}, n_{\text{change}})$$

$$\mathcal{A}_s = \{(1,1) (1,0) (1,-1) (0,1) (0,-1) (-1,1) (-1,0) (-1,-1)\}$$

Reward Function

$$d1 = \frac{|(m_g - m_p)(n_c - n_p) - (n_g - n_p)(m_c - m_p)|}{\sqrt{(m_g - m_p)^2 + (n_g - n_p)^2}}$$

$$d2 = 1 - \frac{|(m_g - m_p)(m_c - m_p) + (n_g - n_p)(n_c - n_p)|}{\sqrt{(m_g - m_p)^2 + (n_g - n_p)^2}}$$

Point 1.  $(m_p, n_p)$  = Previous position reference

Point 2.  $(m_c, n_c)$  = Current position reference

Point 3.  $(m_g, n_g)$  = Goal position reference

$$R = -(1 + d1)(1 + d2)$$

## Training - MATLAB Virtual Robot


- $s = (m_{rel}, n_{rel})$
- $\mathcal{A}_s = \{(1, 1) (1, 0) (1, -1) (0, 1) (0, -1) (-1, 1) (-1, 0) (-1, -1)\}$
- $R = -(1 + d1)(1 + d2)$
- Transition probabilities
  1.  $P(s'|s, a) = 1$  for  $s' = (m_{rel} - m_{change}, n_{rel} - n_{change})$
  2.  $P(s'|s, a) = 0$  for all other  $s'$ .
- $\gamma = 0.95$
- $size(Q) = ((2m_{max} - 1) \times (2n_{max} - 1) \times size(\mathcal{A}))$

```
while( (m_rel ~= 0 || n_rel ~= 0) && (step <= step_max) )  
  
    [max_cur_act, max_cur_index] = max(state_q_values(:, (m_rel + m_max), (n_rel + n_max)));  
    m_rel_next = m_rel + next_rel(max_cur_index, 1);  
    n_rel_next = n_rel + next_rel(max_cur_index, 2);  
  
    if((abs(m_rel_next) > (m_max-1)) || (abs(n_rel_next) > (n_max-1)))  
        state_q_values(max_cur_index, (m_rel + m_max), (n_rel + n_max)) = -Inf;  
    else  
        max_next_act = max(state_q_values(:, (m_rel_next + m_max), (n_rel_next + n_max)));  
        reward = -((m_rel_next)^2 + (n_rel_next)^2)^0.5;  
        state_q_values(max_cur_index, (m_rel + m_max), (n_rel + n_max)) = reward + gamma * max_next_act;  
        m_rel = m_rel_next;  
        n_rel = n_rel_next;  
        step = step + 1;  
    end  
end
```

# ROBOT TASK 1 - NO OBSTACLES; ANY GOAL

## Testing (ARIA) - MobileSim Simulation


```
C:\Program Files\MobileRobot\ARIA\bin\test.exe
Connecting to simulator through tcp.
Syncing 0
Syncing 1
Syncing 2
Connected to robot.
Name: MobileSim
Type: Pioneer
Subtype: p3dx-sh
Loaded robot parameters from p3dx-sh.p
test1: Connected.
Goal Position: < 7159.000000, 6562.000000 >
Goal(n): 15, Goal(n): 14
Current Position: < 0.000000, 0.000000 >
Current(n): 1, Current(n): 1
Relative(n): -14, Relative(n): -13
Policy: 1
Current Position: < 820.000000, 810.000000 >
Current(n): 2, Current(n): 2
Relative(n): -13, Relative(n): -12
Policy: 1
Current Position: < 1182.000000, 1179.000000 >
Current(n): 3, Current(n): 3
Relative(n): -12, Relative(n): -11
Policy: 1
Current Position: < 1780.000000, 1785.000000 >
Current(n): 4, Current(n): 4
Relative(n): -11, Relative(n): -10
Policy: 1
Current Position: < 2207.000000, 2204.000000 >
Current(n): 5, Current(n): 5
Relative(n): -10, Relative(n): -9
Policy: 1
Current Position: < 2812.000000, 2809.000000 >
Current(n): 6, Current(n): 6
Relative(n): -9, Relative(n): -8
Policy: 1
Current Position: < 3174.000000, 3170.000000 >
Current(n): 7, Current(n): 7
Relative(n): -8, Relative(n): -7
Policy: 1
Current Position: < 3816.000000, 3812.000000 >
Current(n): 8, Current(n): 8
Relative(n): -7, Relative(n): -6
Policy: 1
Current Position: < 4137.000000, 4133.000000 >
Current(n): 9, Current(n): 9
Relative(n): -6, Relative(n): -5
Policy: 1
```



# ROBOT TASK 1 - NO OBSTACLES; ANY GOAL

## Testing - P3-DX Physical Robot


Fixed start reference, different goal references



# ROBOT TASK 1 - NO OBSTACLES; ANY GOAL

## Testing - P3-DX Physical Robot



### Fixed goal reference, different start references



# ROBOT TASK 1 - NO OBSTACLES; ANY GOAL

## Testing - P3-DX Physical Robot


### Multiple goals






## Training - MATLAB Virtual Robot

- $s = (m_c, n_c)$
- $\mathcal{A}_s = \{(1, 1) (1, 0) (1, -1) (0, 1) (0, -1) (-1, 1) (-1, 0) (-1, -1)\}$
- $R = -(1 + d1)(1 + d2)$
- Transition probabilities
  1.  $P(s'|s, a) = 1$  for  $s' = (m_c - m_{\text{change}}, n_c - n_{\text{change}})$
  2.  $P(s'|s, a) = 0$  for all other  $s'$ .
- $\gamma = 0.95$
- $\text{size}(Q) = (m_{\text{max}} \times n_{\text{max}} \times \text{size}(\mathcal{A}))$




### Testing (ARIA) - MobileSim Simulation



## ROBOT TASK 2 - FIXED OBSTACLES; FIXED GOAL


## Testing - P3-DX Physical Robot



## Training - MATLAB Virtual Robot

- $s = (m_{rel}, n_{rel}, O_N, O_W, O_E, O_S)$
- $\mathcal{A}_s = \{ (1, 0) (0, 1) (0, -1) (-1, 0) \}$
- $R = -(1 + d_1)(1 + d_2)$
- Transition probabilities
  1.  $P(s'|s, a) = 1$  for  
 $s' = (m_{rel} - m_{change}, n_{rel} - n_{change}, O'_N, O'_W, O'_E, O'_S)$
  2.  $P(s'|s, a) = 0$  for all other  $s'$ .
- $\gamma = 0.95$
- $size(Q) = \left( (2m_{max} - 1) \times (2n_{max} - 1) \times 2 \times 2 \times 2 \times 2 \times size(\mathcal{A}) \right)$

## Testing (ARIA) - MobileSim Simulation



## ROBOT TASK 3 - ANY OBSTACLES; ANY GOAL

### Testing - P3-DX Physical Robot



# CONCLUSION

---

## SUMMARY OF PROJECT

In a problem with discrete state- and action-spaces, by solving three independent tasks of mounting complexity, I successfully trained the robot to go from **any** initial position to **any** goal position, in an unknown environment with **obstacles**, by engaging its SONAR range device. This algorithm may also be applied in dynamic obstacle navigation.

Each task was performed in three stages:

1. Training on a virtual robot using MATLAB.
2. Testing in simulation, using MobileSim, a simulator for MobileRobots. The arena is set up using ARIA, the higher-level software interface for MobileRobots platform.
3. Testing on the physical robot Pioneer 3-DX. <sup>1</sup>

---

<sup>1</sup>Video for testing can be found here - <https://www.dropbox.com/sh/fd2b2p1xi6wsjnu/AACxvR9695aX-UGdan80Ta3ba?dl=0>



1. **Continuous** State and Action Spaces
2. **Options** - An example set of options that may be learnt are:
  - 2.1 Avoid obstacle
  - 2.2 Go to goal
  - 2.3 Follow the wall

- [1] Andrew G Barto. Reinforcement Learning: An Introduction. MIT press, 1998.
- [2] Andrey V Gavrilov and Artem Lenskiy. Mobile robot navigation using reinforcement learning based on neural network with short term memory. In Advanced Intelligent Computing, pages 210–217. Springer, 2012.
- [3] Leslie Pack Kaelbling, Michael L Littman, and Andrew W Moore. Reinforcement learning: A survey. Journal of artificial intelligence research, pages 237–285, 1996.
- [4] Martin Stolle and Doina Precup. Learning options in reinforcement learning. In Abstraction, Reformulation, and Approximation, pages 212–223. Springer, 2002.
- [5] Mu-Chun Su, De-Yuan Huang, Chien-Hsing Chou, and Chen-Chiung Hsieh. A reinforcement-learning approach to robot navigation. In Networking, Sensing and Control, 2004 IEEE International Conference on, volume 1, pages 665–669. IEEE, 2004.

QUESTIONS?